

Hashing trick,
Random projections

Categorical data: 1-hot encoding

	item_type	seller_id
0	book	1234
1	phone	45
2	clothes	45
3	book	46
4	shoes	234
...

$X :=$ "item_type is" 'book',
"seller" is '1234'

$X \in \mathbb{R}^d$

- List all modalities of "item_type"
- Assign one index to each modality.
- Idem for other variables

```
{ 'shoes': 1,  
  'phone': 2,  
  'book': 3,  
  'clothes': 4,  
  ... }
```

0 type is "shoes"
0 type is "phone"
1 type is "book" ✓
0 type is "clothes"
...
0 seller is "45"
1 seller is "1234" ✓
0 seller is "5678"
...

$\in \mathbb{R}^d$ with
 $d :=$ nb distinct "types" + nb distinct "seller"

Polynomial Kernel and « Crossfeatures »

“Quadratic Kernel”

Formally: $X.X^T \in \text{Matrices}^{d,d} \cong \mathbb{R}^{d^2}$

Or use index:

```
('phone', '45'): 1000,  
( 'phone', '46'): 1001,  
( 'phone', '47'): 1002,  
...  
( 'book', '45'): 1013,  
( 'book', '46'): 1014,  
...
```

			0 0 1 0 ... 0 1 0 ...	X^T
0	type is “shoes”	0	0 0 0 0 ... 0 0 0 ...	
0	type is “phone”	0	0 0 0 0 ... 0 0 0 ...	
1	type is “book” ✓	1	0 0 1 0 ... 0 1 0 ...	
0	type is “clothes”	0	0 0 0 0 ... 0 0 0 ...	
...		
0	seller is “45”	0		
1	seller is “1234” ✓	1		
0	seller is “5678”	0		
...		...		

$X \in \mathbb{R}^d$

- Type is “book” and seller is “1234”
“Crossfeature” between “type” and “seller”

```
0 type is “shoes”  
0 type is “phone”  
1 type is “book” ✓  
0 type is “clothes”  
...  
0 seller is “45”  
1 seller is “1234” ✓  
0 seller is “5678”  
...  
0 “book” and “45”  
1 “book” and “1234” ✓  
0 “book” and “5678”  
...
```

$\in \mathbb{R}^{d^2}$

Help I now have too many features !

- I have quite a few high cardinality categorical columns
 - Even building a vocabulary of possible values is a pain
- I want to use cross features and I have too many of them
- I want to limit the size of my model

Hash functions

a hash function:

- Maps complex data to integers
- Deterministic
- But « looks random »

```
1 print(hash("a"))
2 print(hash("b"))
3 print(hash("hello"))
4 print(hash("this is a longer string"))
5 complex_object = ("some_structured_data", 123, 45)
6 print(hash( complex_object ))
7 print(hash("a"))
```

```
4622726560737668235
-6279647101590938825
2716929806936918215
-6349598305679374822
6867566546198076494
4622726560737668235
```

Hash("a")
Deterministic!

- Usages:

- Pseudo random numbers
- Partitioning data
- Internal implementation of « dictionary »
- ... *surprisingly useful!*

```
1 # Processing ~= 1% of items:
2 for item in myList:
3     if hash(item) % 100 == 0:
4         process(item)
5 # If the same item appears twice, it will be either processed twice, or not at all.
6 # If we run again, the same items will be selected, even if list is shuffled
```

The Hashing Trick

Building the index *without* looking at the data

- *Hash everything!*
- *Write “1” at the indexes of the hashes*

“type” is “book”,
“seller” is “1234”

Raw input

```
hash( ("type", "book")) %1000000
```

966886

```
hash( ("seller", "1234")) %1000000
```

175421

```
hash( (("seller", "1234"), ("type", "book"))) %1000000
```

548128

Chosen size d
of the output
vector

1 at index
Hash(seller, 1234)

0
0
..
0
1
0
..
0
1
0
..
0
1
0
..

Hash of the
“crossfeature”

$X \in \mathbb{R}^d$ with
 $d := 1000000$
(your choice)

Collisions?

1	<code>hash(("seller", "626"))%1000000</code>
executed in 7ms, finished 15:19:59 2021-03-09	

626786

1	<code>hash(("type", "1674"))%1000000</code>
executed in 4ms, finished 15:19:59 2021-03-09	

626786

0
0
...
0
1
0
...

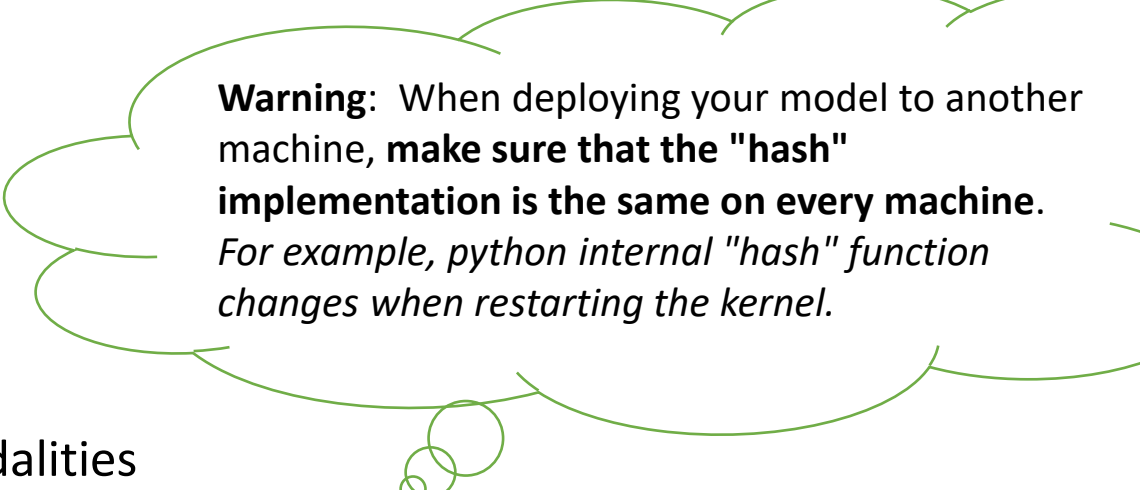
A "1" at index 626786 may either mean:

- 'type is 1674'
- or 'seller is 62'
- Or something else ? (more hashes colliding)

- When learning a model from hashed inputs (eg a linear model $X.w$), collisions constrain the learned model to set the **same weights to colliding features**.
- Here, same weight for «seller is 626» and «type is 1674»

Hashing Trick: Why?

- Simple
 - No need to precompute indexes
 - No need to care *too much* about « new » modalities
- Fixed size model
 - **Known memory footprint** → no surprise crash
- Collisions are not so bad
 - Collision between rare modalities is OK
 - Collisions between frequent modalities is unlikely
 - Also, redondant features may help the model to recover from collisions (eg lots of « crossfeatures »)
- Common use cases
 - (Generalized) linear model with many crossfeatures
 - Input layer of neural networks



Warning: When deploying your model to another machine, **make sure that the "hash" implementation is the same on every machine.** For example, python internal "hash" function changes when restarting the kernel.

Avoiding Collisions?

```
1 hash(( "seller", "626" ))%100000000
```

12626786

```
1 hash(( "type", "1674" ))%100000000
```


86626786

- Increase $n \rightarrow$ less collisions.
- Difficult to fully avoid: « birthday paradox »
 - K distinct hashed modalities $\Rightarrow O(K^2)$ pairs of modalities
 - Each single pair has a probability $1/n$ of colliding together
 - Overall, $O(K^2/n)$ expected collisions
- But increasing n increases the size of the model and the learning time
- Test for best tradeoff!!

- Other options:
 - Keep index of «common» modalities and hash only «rare» modalities
 - Several hashes per modality (more collisions, but more redundancy to recover from those collisions)
- But try «Ostrich» algorithm first!!

The hashing trick


<https://booking.ai/dont-be-tricked-by-the-hashing-trick-192a6aae3087>



Booking dataset

The hashing trick


<https://booking.ai/dont-be-tricked-by-the-hashing-trick-192a6aae3087>



Criteo dataset


The hashing trick

<https://booking.ai/dont-be-tricked-by-the-hashing-trick-192a6aae3087>



Avazu dataset

Dimensionality reduction



Reducing the dimensions of input data before applying ML algo.


Ideally:

- $K \ll D$
- (approximately) preserving distances between examples:
 $d(X_i, X_j) \approx d(f(X_i), f(X_j))$
- f linear. $f: X_i \rightarrow A.X_i$. (easy to compute)

Classical method: PCA

- Rows of $A :=$ first K eigenvectors of $X \cdot X^T$
- Good preservation of distances ✓
- $O(ND^2)$ or $O(N^2D)$ ✗
- Data dependent
(can we avoid it ?)

Random projections



- A is a matrix of size (K,D)
- Crazy idea: **Choosing A randomly!**?

What if...

- I Choose $A_{ij} \in \{0;1\}$
- Independent columns, with exactly one 1 per column, placed uniformly at random ?

This is hashing trick!! (if X is binary)

- Still requires large K ($> 10^6$) to work well

Instead, “Gaussian” random projections:

$$A_{ik} \text{ iid } \sim \text{Gaussian}(0,1) * 1/\sqrt{K}$$

- Is it likely to get $d(A.X_i, A.X_j) \approx d(X_i, X_j)$?
- For all pairs i,j simultaneously ??

Random projections

Let $z \in \mathbb{R}^D$ and $|| \cdot ||$ the Euclidian norm.

Let's compare $|A.z|^2$ and $|z|^2$.

$$|A.z|^2 = \sum_k (A_k \cdot z)^2$$

Lemma:

- $A_k \cdot z$ are iid
- $A_k \cdot z \sim N(0,1) * |z| / \sqrt{K}$

Thus:

- $E((A_k \cdot z)^2) = |z|^2 / K$
- $E(|A.z|^2) = |z|^2$
- $\text{Var}(|A.z|^2) = 2|z|^4 / K$

Checking that " $|A.z| \approx |z|$ " (with high proba)

Set $z := X_i - X_j$ to get

$$d(X_i, X_j) \approx d(A.X_i, A.X_j)$$

A sum of independent Gaussian variables is Gaussian:

$$A_k \cdot z = \sum_j A_{k,j} * z_j$$

$A_{k,j} * z_j$ are independent and $\sim N(0, |z_j| / \sqrt{K})$

(Variance of a χ^2)

$O(1/K)$: With large K , low variance.

Large probability of getting $|A.z| \approx |z|$

Random projections

Let $\varepsilon > 0$. We would like, for all pair (i,j) :
 $(1 - \varepsilon) \cdot d(X_i, X_j) < d(A.X_i, A.X_j) < (1 + \varepsilon) \cdot d(X_i, X_j)$

Formalizing

$$"d(X_i, X_j) \approx d(A.X_i, A.X_j)"$$

No 100% garanty that inequalities are true.
Instead, we want them to hold (for all i,j)
with a probability $> 1 - \delta$

Let $\delta > 0$ « *accepted Probability of failing* »

Theorem (Johnson–Lindenstrauss):

if $K > O(\log(n/\delta) / \varepsilon^2)$, the probability that all inequalities above hold is higher than $1 - \delta$

- Does not depend on D !
- $O()$ constant is around 10

Proof idea: $O(n^2)$ pairs \rightarrow Let us get Proba of failing for one pair smaller than δ/n^2
Use concentration of $|Az|^2$ around its mean.

- Example:
 $N=1B, \varepsilon = 0.1 \rightarrow K \approx 20000$
- In practice, K significantly smaller is often good enough.

Random projections

Typical values

$n = O(1B); d = O(1M) \Rightarrow k = O(10K)$

Memory: storing A : $O(DK) \approx 40GB$

- Maybe ok?
- Use only a few bits per entry?
- Store only the seeds?!

Computing $A.X$: $O(ndk)$ **X**

- But X is usually sparse **✓**

Many variants

- Replace $A_{ij} \sim \text{Gaussian}$ by A_{ij} in $\{+/-1\}$
 - 1 bit per entry
- "Fast Johnson-Lindenstrauss transform"
 - A is the product of sparse random matrices and a structured matrix allowing fast computation
 - $O(n.d.\log(d) + nk^2)$ **✓**