

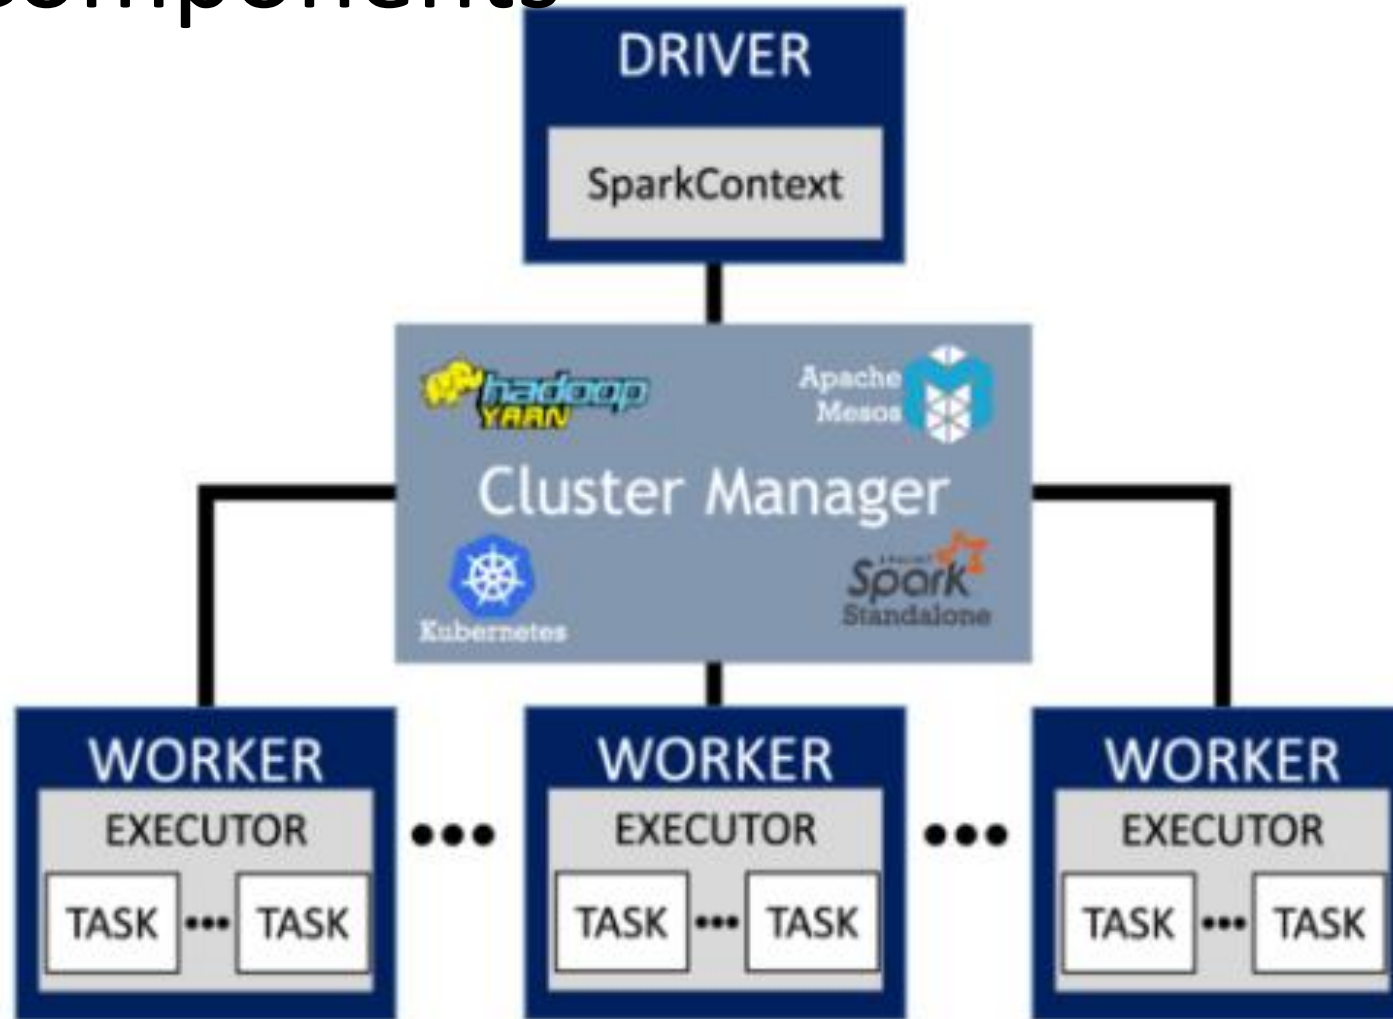
Systems, paradigms and algorithms for Big Data

TD 1

Why Spark ?

- Ease creation of complicated data analysis task
- Support iterative algorithms, like gradient descent
- Ease data exploration
- Make it easier to explore data interactively
 - Developer friendly, no need to create tons of classes and jobs...
 - Repl mode
 - Scala/Python API similar to functional programming

Main Components



How to run it

- Repl : Spark-shell, notebooks → exploratory mode
- Spark-submit : runs a script → scheduled job
- Api for Java, Scala, Python.
- Deploy Mode : client vs cluster

Dependencies

- Depends on Hadoop Libraries (HDFS and Yarn)
- Requires a Java Runtime Environment (need Java in your system path)

RDD

Resilient Distributed Dataset

- The base building block of your application
- Lazy : doesn't compute anything before it really needs to do so.
- Immutable : a transformation doesn't change the data set, it returns a new RDD.
- Fault Tolerant : partition can be recomputed in case of failure

Graph

- Conceptually, an RDD is a node of a graph of functions.
- API is pretty much similar with **Functional Programming. Forget the for-loop.**

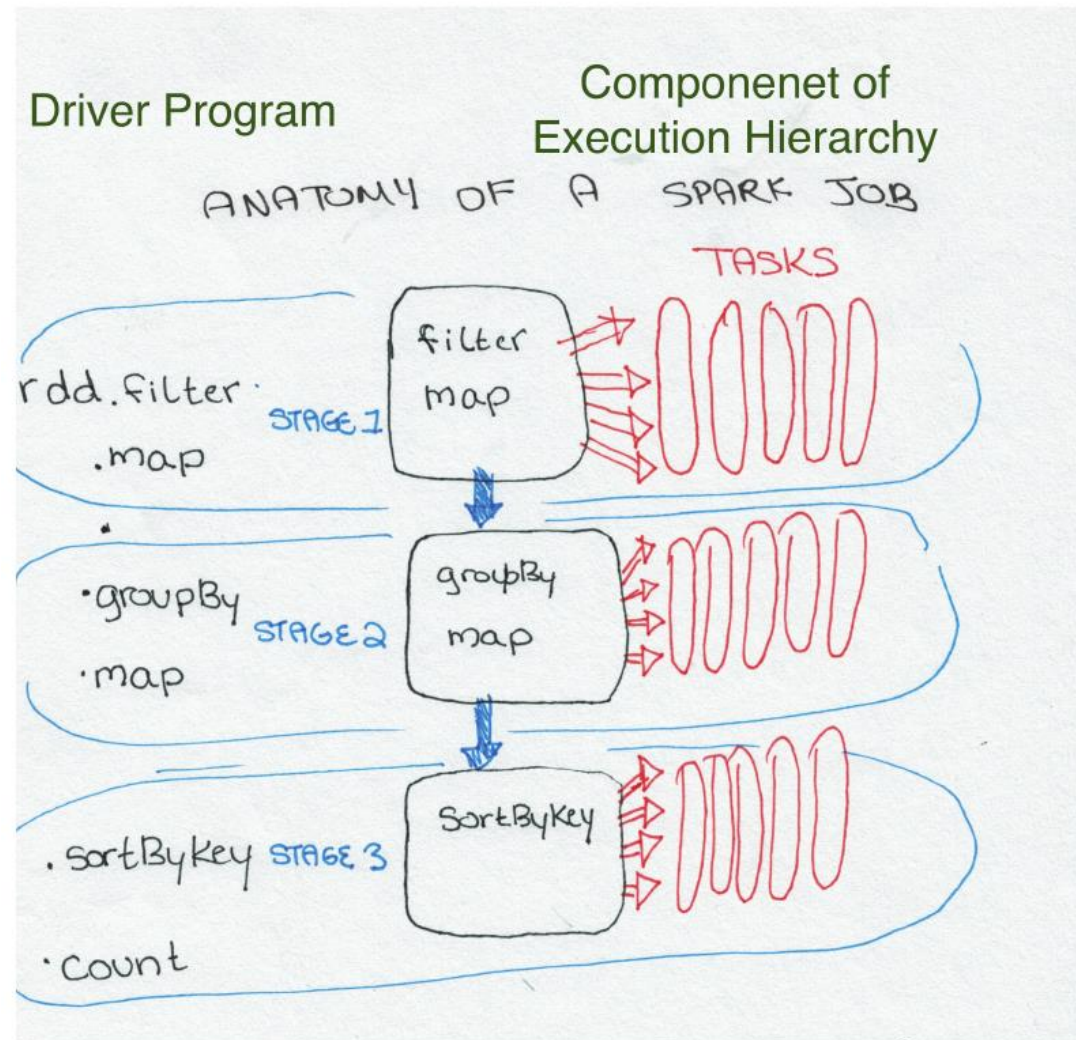
Inside the RDD:

- Knows the partitions he is working-on, how data is partitionned
- Knows how to iterate over each partition to yield records
- Know RDD's it depends-on

RDD as a Directed Acyclic Graph




Jobs, Stages, Tasks




Source : high performance Spark

RDD API - Actions


Row
{"user":"John", "movie":"Blade Runner", "rating":5.0}
{"user":"Louise", "movie":"Dirty dancing", "rating":5.0}
{"user":"Sam", "movie":"Blade Runner", "rating":3.5}



Row
{"user":"John"}
{"user":"Louise"}
{"user":"Sam"}



Row
{"user":"John"}
....
{"user":"Sam"}



RDD API

Row

```
{"user":"John", "movie":"Blade Runner", "rating":5.0}  
{"user":"Louise", "movie":"Dirty dancing", "rating":5.0}  
{"user":"Sam", "movie":"Blade Runner", "rating":3.5}
```

Row

```
{"user":"John", "movie":"Blade Runner", "rating":5.0}  
{"user":"Louise", "movie":"Dirty dancing", "rating":5.0}  
{"user":"Sam", "movie":"Blade Runner", "rating":3.5}
```

Row

```
("John", "Blade Runner")  
("Louise", "Dirty dancing")  
("Sam", "Blade Runner")
```

map($x \rightarrow x['movie']$)

Row

```
Blade Runner  
Dirty Dancing  
Blade Runner
```

keyBy($x \rightarrow x['user']$)

Row

```
("John", ...)  
("Louise", ...)  
("Sam", ...)
```

mapValues($x \rightarrow$
 $\text{len}(x[1])$)

Row

```
("John", 12)  
("Louise", 13)  
("Sam", 12)
```

RDD API

Row
{"movie": "Blade Runner", "genres": "cyberpunk;scifi;action"}
{"movie": "Dirty dancing", "genres": "music;dance;romance"}

flatMap(x →
x['genres'].split(';'))

Row

cyberpunk
scifi
action
music
danse
romance

Row
{"user": "John", "movie": "Blade Runner", "rating": 5.0}
{"user": "Louise", "movie": "Dirty dancing", "rating": 5.0}
{"user": "Sam", "movie": "Blade Runner", "rating": 3.5}


filter(x →
x['rating'] > 4.0)

Row

{"user": "John", "movie": "Bl...
{"user": "Louise", "movie": "Dir...


RDD API - Aggregations

Row
5.0
5.0
3.5



Row
13.5

Row
("Blade Runner", 5.0)
("Dirty dancing", 5.0)
("Blade Runner", 3.5)



Row
("Blade Runner", 8.5)
("Dirty dancing", 5.0)

Other useful functions

- **Join (shuffle ? Lazy ?)**
- **Sample**
- **mappartitions**
- **zippartitions**

Links

- <https://spark.apache.org/docs/latest/api/scala/org/apache/spark/api/java/JavaPairRDD.html>
- <https://0x0fff.com/hadoop-mapreduce-comprehensive-description/>